

Adaptive Trust-Oriented Runtime Security for AI-Generated Shell Scripts

Denys Yakymov

ORCID: 0009-0005-2398-8976 · yakden@gmail.com

Abstract—Large language models now write shell scripts that reach production with little review, and a script that lints cleanly can still wipe a disk, open a reverse shell, or ship credentials off the host. We present ACTE (Adaptive Context-aware Trust Execution), a pre-execution trust gate that scores a script through a transparent thirteen-feature logistic risk model over semantic, execution-context, and threat-signature evidence, and compiles the verdict into an enforcement policy. We lead with the honest generalisation numbers. On a reproducible corpus of 420 labelled scripts, five-fold cross-validation gives $F1 = 0.887 \pm 0.025$ at the default operating point; a stricter leave-template-out protocol that forbids any template from spanning train and test yields $F1 = 0.823 \pm 0.101$, and cluster (block) bootstrap intervals that respect the template structure are about twice as wide as naive ones ($F1$ in $[0.79, 0.99]$). Recall on the genuinely hard malicious class is 0.75, not the 0.855 aggregate, and the false-positive rate at the default threshold is ≈ 0.10 , not the 0.011 of a single tuned split. Trained only on synthetic data and frozen, ACTE clears every safe script in an independent 41-script real holdout (0 of 20 false positives, Wilson 95% CI $[0, 0.16]$) but misses evasive payloads its signatures do not cover (recall 0.81). Against strong learned baselines—TF-IDF logistic regression, a linear SVM, a random forest, a gradient-boosted tree on the same features, and a feature+TF-IDF union—ACTE is competitive, not dominant, on raw $F1$ (a GBDT on its own features matches it; the union reaches 0.99 on this near-separable corpus). Its value is elsewhere: the lowest false-positive rate, an auditable decision, and a really-enforced policy—on Linux we compile a content-derived syscall deny-set into an actual seccomp-BPF filter and show the kernel enforce it end-to-end. Under adaptive evasion the two model families show complementary blind spots: additive benign camouflage collapses the lexical baseline but not ACTE, whereas behaviour-preserving semantic substitution evades ACTE (recall $0.86 \rightarrow 0.58$) but not the lexical baseline, so neither is robustly superior. We are candid about the limits—the training corpus is synthetic and author-generated, the score is only moderately calibrated (ECE 0.076, 0.055 after isotonic recalibration), the online-learning loop shows no demonstrated benefit and is poisonable with a handful of mislabelled updates, and there is no independent human labelling. We position ACTE as a reproducible, honestly-evaluated framework—interpretable content-to-policy trust gating with real enforcement—rather than a finished detector.

Index Terms—runtime security, shell-script analysis, AI-generated code, trust evaluation, logistic risk model, online learning, seccomp, Open Policy Agent, zero-trust execution.

I. INTRODUCTION

The proliferation of large language models capable of generating executable code has fundamentally altered the software supply-chain landscape. Tools such as GitHub Copilot, ChatGPT, and their successors are now routinely employed to produce shell scripts for infrastructure management, deployment automation, and one-off administrative tasks. While this accelerates development velocity, it introduces a category of risk that conventional static analysis tools were not designed to address: the AI-generated shell script as an attack surface.

Shell scripts are privileged, imperative artifacts. A single malicious or misconfigured script can exfiltrate credentials, establish persistence, destroy filesystems, or exfiltrate data over covert channels. The key threat model concern is not that LLMs are inherently malicious, but rather that

(i) prompt injection and adversarial prompting can induce dangerous outputs, (ii) model hallucination can produce syntactically valid yet semantically harmful constructs, and (iii) the opacity of LLM generation means that standard code-review heuristics fail to capture the full behavioral envelope of the generated code.

Existing defenses are inadequate for this threat model. Static linters such as ShellCheck [1] detect syntactic errors and style violations but are agnostic to intent and behavioral context. Mandatory access-control frameworks such as SELinux [2] and AppArmor [3] operate at the process level and require pre-authored policy that cannot adapt to unseen code at runtime. Behavioral monitors such as Falco [4] and eBPF-based auditing systems [5] observe execution but cannot preempt it. None of these approaches integrates a continuous risk score derived from the script's content,

execution context, and historical feedback into a unified adaptive trust decision made *before* the script runs.

We present the **Adaptive Context-Aware Trust Execution (ACTE)** framework to address this gap. ACTE's central contribution is a formal logistic risk function $R(s)$ that aggregates features from three orthogonal evidence dimensions—semantic structure, execution context, and threat-intelligence signatures—into a continuous risk score in $(0, 1)$, which is inverted to a trust score $T(s) = 1 - R(s)$ and mapped to a discrete enforcement level. The weight vector is updated online by SGD whenever labelled feedback arrives, so the decision boundary can drift toward a given deployment's script distribution without a batch retrain.

The paper contributes four things. We formalise the problem of pre-execution trust scoring for AI-generated shell scripts, with an explicit threat model (Section II), and give a compact logistic risk model with an online update rule and a documented monotonicity property (Sections III–IV). We describe an eight-module architecture that turns a trust score into concrete enforcement artefacts, and a Python prototype (Sections V–VI). We evaluate that prototype under a deliberately demanding protocol—cross-validation, a leave-template-out split, bootstrap intervals, learned baselines, and an independent real-world holdout—and report the results, favourable and unfavourable alike (Sections VII–IX). And we are equally explicit about what the work does not establish (Section X): the training data is synthetic and author-generated, enforcement is simulated, and there is no adaptive-attacker evaluation. We see the value of the paper as a reproducible framework and an honest baseline, not as a finished detector.

II. PROBLEM FORMALIZATION

A. Object of Study

Let $\mathcal{S}$ denote the universe of POSIX-compatible shell scripts. Each script $s \in \mathcal{S}$ is a finite sequence of characters that may be submitted to a shell interpreter for execution on a host system with resources $\mathcal{R}$ (filesystems, network interfaces, device nodes, process table). We are interested in the subset $\mathcal{S}_{\text{AI}} \subseteq \mathcal{S}$ of scripts produced, in whole or in part, by generative language models, though the framework is applicable to any shell script regardless of provenance.

B. Threat Model

We consider an adversary $\mathcal{A}$ who controls the prompt submitted to an LLM-based code-generation service and whose goal is to cause the generated script to perform one or

more of the following: (i) unauthorized data exfiltration, (ii) persistence establishment, (iii) privilege escalation, (iv) irreversible data destruction, or (v) lateral movement. We assume the adversary can craft obfuscated scripts (e.g., base64-encoded payloads, eval-based indirection, character-substitution ciphers) to evade pattern-based defenses. We do not model adversaries with physical hardware access or kernel-level root kits; our threat boundary ends at the system-call interface.

C. Formal Definitions

Let $\mathcal{C}$ be the set of distinct commands (executables) that may appear in scripts in $\mathcal{S}$. Let $\mathcal{X} : \mathcal{S} \rightarrow [0, 1]^d$ be a feature-extraction function mapping any script to a d -dimensional real-valued feature vector. Define a *risk function* $R : \mathcal{S} \rightarrow (0, 1)$ and a *trust function* $T : \mathcal{S} \rightarrow (0, 1)$ such that $T(s) = 1 - R(s)$ for all s . The *trust level* of a script is a discrete label $\ell(s) \in \{\text{TRUSTED}, \text{MONITOR}, \text{RESTRICT}, \text{DENY}\}$ obtained by thresholding $R(s)$.

D. Research Questions

We structure the experimental evaluation around four research questions (RQs):

RQ1 (Detection Accuracy): What classification performance does ACTE achieve on held-out labeled scripts; how stable is that performance under cross-validation; and how does each component contribute?

RQ2 (Computational Cost): Is ACTE's per-script analysis latency compatible with inline deployment on interactive or CI/CD pipelines?

RQ3 (False-Positive Reduction): By how much does ACTE reduce the false-positive rate relative to a ShellCheck v0.9.0 baseline, and is the difference statistically significant?

RQ4 (Generalization to Real Scripts): Does a model trained solely on the synthetic corpus generalize to independent, real, publicly-documented scripts it has never seen?

III. THE ACTE MODEL AND NOVELTY

We are careful about what is and is not novel. Operating before execution and emitting a continuous score are properties ACTE shares with every static analyser, so we do not claim them as contributions. What the framework offers as a package is threefold. First, it derives a per-script enforcement policy from the script's *content*—not merely from a risk tier—and, on Linux, really enforces it: a content-derived syscall deny-set is compiled to a seccomp-BPF filter and installed, so the trust decision becomes an installed

kernel control rather than a document (Section VIII-L). Second, every decision decomposes into thirteen named, signed indicators; we argue this is a security property and not merely an ergonomic one, though the adversarial results of Section VIII-H qualify how far it goes. Third, the whole content $\rightarrow$ risk $\rightarrow$ artefact pipeline is packaged and evaluated end to end against a concrete threat. Online adaptation is part of the design, but we treat it as an unproven mechanism rather than a contribution, because our experiments (Section VIII-K) show no net benefit and expose a poisoning surface.

The four trust levels map to concrete enforcement postures: TRUSTED scripts execute without restriction; MONITOR scripts execute with full audit logging; RESTRICT scripts execute in a sandboxed namespace with network egress blocked, read-only root filesystem, and capped CPU/memory/PID resources; DENY scripts are blocked entirely with no execution permitted.

The online feedback loop closes the adaptation loop: when a MONITOR-level script is later confirmed safe or dangerous by an operator, a single SGD step updates the weight vector in sub-millisecond time, moving the decision boundary for future similar scripts without requiring batch retraining.

IV. FORMAL MATHEMATICAL MODEL

A. Feature Space and Sets

Let $\mathcal{C}$ be the command set extracted by the SemanticParser, and let $\mathcal{G} = \{\text{semantic}, \text{context}, \text{threat}\}$ be the three feature groups. The full feature vector is $\mathbf{x}(s) \in [0, 1]^{13}$, with 5 semantic features, 5 context features, and 3 threat-intelligence aggregates (see Section V for definitions).

B. Risk Function

The logit (linear combination) is

$$z(s) = b + \sum_{i=1}^{13} w_i x_i(s),$$

where b is a scalar bias and $\mathbf{w} \in \mathbb{R}^{13}$ is the weight vector.

The risk score is the logistic (sigmoid) transformation:

$$R(s) = \sigma(z(s)) = \frac{1}{1 + e^{-z(s)}} \in (0, 1).$$

C. Trust Function and Levels

Trust is the complementary quantity:

$$T(s) = 1 - R(s) \in (0, 1).$$

The discrete trust level is determined by

$$\ell(s) = \begin{cases} \text{TRUSTED} & R(s) < 0.25, \\ \text{MONITOR} & 0.25 \leq R(s) < 0.50, \\ \text{RESTRICT} & 0.50 \leq R(s) < 0.80, \\ \text{DENY} & R(s) \geq 0.80. \end{cases}$$

Binary detection for evaluation purposes uses a tunable decision threshold τ (default 0.50): a script is predicted dangerous if $R(s) \geq \tau$.

D. Monotonicity Property

We present this property but caution (Sections VIII-H, X-D) that in a security setting it is a transparency guarantee with an adversarial cost, not a robustness guarantee: it blocks additive evasion but not semantic substitution [23], [26]. Because σ is strictly increasing, $R(s)$ is monotone in $z(s)$: adding a positive-weighted feature never decreases risk, and adding a benign-weighted (negative) feature never increases risk. Formally, if $\mathbf{x}'(s) \geq \mathbf{x}(s)$ component-wise (with at least one strict inequality and $w_i > 0$ for the differing components), then $R(s') > R(s)$. This property provides an auditable safety invariant useful for formal analysis and operator trust.

E. Online Trust-Update Algorithm

When a labeled example $(s, y) \in \mathcal{S} \times \{0, 1\}$ is received (with $y = 1$ indicating dangerous), FeedbackLearning applies a single SGD step on the binary cross-entropy loss $\mathcal{L}(p, y) = -[y \ln p + (1 - y) \ln(1 - p)]$, where $p = R(s)$:

Algorithm 1: Online SGD Weight Update

```

Input: features  $\mathbf{x}$ , label  $y \in \{0, 1\}$ , learning
rate  $\eta$ , L2 penalty  $\lambda$ 
 $z \leftarrow b + \sum_i w_i \cdot x_i$ 
 $p \leftarrow \sigma(z)$ 
error  $\leftarrow p - y$  // gradient of
BCE wrt logit
for each feature  $i$ :
    grad  $\leftarrow$  error  $\cdot x_i + \lambda \cdot w_i$  // gradient +
    L2 regularization
     $w_i \leftarrow w_i - \eta \cdot \text{grad}$ 
 $b \leftarrow b - \eta \cdot \text{error}$ 
Output: updated ( $\mathbf{w}$ ,  $b$ )

```

The L2 penalty $\lambda \|\mathbf{w}\|^2$ prevents weight drift from auditable hand-set priors. Across 40 training epochs on 252 labeled scripts, the weight vector shifted by $\|\Delta \mathbf{w}\| = 6.08$ and loss declined from 0.345 to 0.288, confirming non-trivial adaptation.

V. SYSTEM ARCHITECTURE

ACTE comprises eight cooperating modules arranged as a linear analysis pipeline followed by an enforcement layer: a

raw script flows through the SemanticParser, ContextExtractor, and ThreatIntel into the TrustEvaluationEngine, whose assessment feeds the PolicyGenerator and RuntimeMonitor, with FeedbackLearning updating the engine's weights out of band.

A. SemanticParser

The SemanticParser converts a raw script string into a structured `ParsedScript` bundle. It attempts to parse each line with `bashlex`, a Python port of the GNU Bash parser, and falls back to a robust `shlex/regex` tokenizer when `bashlex` raises (as it frequently does on obfuscated or adversarial input). It extracts: command names, flags, file paths, URLs, device nodes, pipe count, pipe-to-shell detection, command-substitution density, redirection count, background-job count, and heredoc count.

B. ContextExtractor

The ContextExtractor derives five normalized execution-context dimensions from the `ParsedScript`: (1) *privilege_required*—presence of privilege-escalation commands (`sudo`, `su`, `setcap`) or sudoers manipulation; (2) *network_egress*—network-capable commands, URLs, pipe-to-shell patterns, or `/dev/tcp` redirections; (3) *filesystem_scope*—access to sensitive roots (`/etc`, `/boot`, `/dev`) or recursive wildcard operations; (4) *irreversibility*—destructive commands (`rm -rf`, `dd`, `mkfs`, `shred`) or no-preserve-root patterns; (5) *obfuscation*—base64 decoding, eval, hex-escape sequences, or high command-substitution density.

C. ThreatIntel

ThreatIntel loads a JSON knowledge base of regex-pattern signatures, each assigned a signed weight indicating threat (>0) or benign (<0) signal strength. It produces a list of `SignatureHit` objects, from which three aggregate features are derived: total weight norm, maximum single-category weight norm, and count of distinct threat categories (each normalized to $[0,1]$).

D. TrustEvaluationEngine

The TrustEvaluationEngine assembles the 13-dimensional feature vector from the outputs of the three preceding modules, evaluates the logistic risk function $R(s)$, and returns a `TrustAssessment` containing the risk score, trust score, trust level, logit, per-feature contributions, and matched signatures.

E. FeedbackLearning

FeedbackLearning exposes an `update_one` method (single SGD step) and a `train` method (multi-epoch batch training over labeled pairs). It additionally provides a `tune_threshold` method that selects the decision threshold τ maximizing F1 (or minimizing FPR subject to a recall constraint) on a validation set.

F. PolicyGenerator

PolicyGenerator translates a `TrustAssessment` into an `ExecutionPolicy` consisting of a seccomp profile sketch (Docker JSON format), an OPA Rego policy snippet, and a namespace/cgroup constraint bundle. Each trust level yields a distinct syscall allow-list, network posture, resource quotas, and capability drop set.

G. RuntimeMonitor

RuntimeMonitor provides a simulation of seccomp-bpf enforcement against a stream of observed syscall names. Given a policy, it builds allow/audit/block/kill decision tables and can be queried per-syscall or in batch mode. In the current prototype, enforcement is simulated (no kernel-level seccomp filter is installed); the monitor is included to demonstrate the enforcement logic in a testable form.

H. ACTEPipeline

The ACTEPipeline orchestrates all seven preceding modules into a single `analyze(script)` call that returns a fully populated `ACTEResult` object, including the latency measured via `time.perf_counter()`.

VI. PROTOTYPE IMPLEMENTATION

The ACTE prototype is implemented in Python 3.11. The SemanticParser uses the `bashlex` library for AST-level parsing and the standard `shlex` module for the fallback tokenizer. Regex-based extraction of URLs, device nodes, and file paths is handled by compiled patterns applied to the raw script string, providing resilience against obfuscated syntax that defeats AST parsers.

The TrustEvaluationEngine's logistic model is implemented entirely in pure Python using the standard `math` module, with a numerically stable sigmoid that avoids overflow for large positive or negative logits. Default weights were hand-set to produce a reasonable cold-start model and are subsequently adapted by FeedbackLearning.

The FeedbackLearning module implements SGD with L2 regularization. Epoch shuffling is performed with a deterministic linear-congruential generator (LCG) seeded from the training invocation, ensuring reproducible training

order independent of any global random state. The threshold tuner performs a sweep over all candidate thresholds present in the scored validation set.

Policy generation produces seccomp JSON in the Docker seccomp profile format and Open Policy Agent Rego snippets. The Rego output specifies `allow`, `audit`, `network_egress_allowed`, and `allow_new_privileges` decisions for each trust level. Namespace constraints specify Linux namespaces to isolate (PID, mount, UTS, IPC, and optionally NET), cgroup CPU quota, memory limit, and PID limit. These artifacts are emitted as data structures; in a production deployment they would be passed to a container runtime or an eBPF enforcement agent.

The complete codebase is released publicly (see the Reproducibility section) and includes the dataset generator, the real-world holdout builder, the experiment runner, the cross-validation and statistics modules, the figure-generation scripts, and the full test suite. The run command `python -m experiments.run_all` reproduces all reported numbers end-to-end.

VII. EXPERIMENTAL METHODOLOGY

A. Dataset

We constructed a dataset of 420 labeled shell scripts spanning five categories: *ai_generated* (scripts plausibly attributable to LLM generation), *malicious* (known-bad payloads), *obfuscated* (scripts with encoding/eval indirection), *safe_everyday* (routine user scripts with no dangerous behavior), and *sysadmin* (legitimate administrative scripts, some with moderate privilege use). The dataset is synthetic-but-reproducible: all scripts were generated by a deterministic dataset generator seeded at 1337, ensuring any researcher can reconstruct the exact dataset and reproduce reported numbers. Hard negatives (safe scripts resembling dangerous patterns) and hard positives (dangerous scripts with minimal syntactic markers) are included to stress-test the classifier. The dataset is split 60/40 train/test (252/168 scripts), yielding 114 positive training examples and 76 positive test examples.

B. Hypotheses

H1 (RQ1): ACTE achieves $F1 > 0.85$ on the held-out test set, and the cross-validated mean F1 remains above 0.80. H2 (RQ2): Mean per-script latency is below 2 ms. H3 (RQ3): ACTE's false-positive rate is at least 80% lower than ShellCheck v0.9.0, and the difference is statistically significant ($\alpha = 0.05$). H4 (RQ4): On the real-world holdout, ACTE retains $F1 > 0.85$ with $FPR \leq 0.10$.

C. Measurement Methodology

For RQ1, ACTE is trained for 40 epochs on the training split. The decision threshold is chosen on the *training* split alone, by maximising F1; for the full model this procedure selects $\tau = 0.50$, and we use that operating point on the held-out test split. (When the model is retrained on a different corpus the tuned threshold moves accordingly; for the full-corpus model of RQ4 it is $\tau = 0.32$.) Classification metrics—precision, recall, F1, MCC, accuracy, ROC-AUC, PR-AUC, FPR, FNR—are then computed on the held-out test split. The ablation study removes one component at a time (SemanticParser, ContextExtractor, ThreatIntel, FeedbackLearning) by zeroing the corresponding feature group and re-running the identical training and evaluation protocol. We caution that the four defined trust levels use fixed *risk* thresholds (DENY at $R \geq 0.80$), which are a separate concern from the binary decision threshold τ used for detection metrics; a script at RESTRICT ($R \geq 0.50$) counts as "flagged" under $\tau = 0.50$.

For RQ2, each of the 420 scripts is analyzed three times and the minimum latency is recorded (to minimize scheduling jitter). Throughput is computed as the reciprocal of the mean latency.

For RQ3, ShellCheck v0.9.0 is applied to each test script; a script is predicted dangerous when any finding at severity `error` or `warning` is produced. We stress at the outset what this comparison can and cannot show. ShellCheck is a correctness and portability linter, not a security classifier, so beating it does not establish state-of-the-art detection; it quantifies the gap that motivates the work—that the tool teams already run is blind to the behaviours we care about. Because both detectors see the same items, we test the difference with the exact binomial form of McNemar's test on the discordant pairs [17], the appropriate paired test for two classifiers on one set. The more demanding comparison, against learned classifiers, is RQ5.

For RQ5, we train three off-the-shelf supervised baselines on TF-IDF features of the raw script text (word 1–2 grams and character 3–5 grams): logistic regression, a linear SVM, and a random forest, all from scikit-learn [16]. Each is evaluated on the same test split as ACTE and, after retraining on the full synthetic corpus, on the real-world holdout, so the transfer comparison is like-for-like.

To move past a single split we cross-validate the full model two ways using scikit-learn's stratified splitters [16].

Stratified five-fold preserves the label ratio in each fold and estimates generalisation to fresh samples from the same distribution. *Leave-template-out five-fold* (a stratified group split keyed on the generating template, of which there are 70) keeps every template wholly inside one fold, so the model is always scored on template structures absent from its training folds; this is the harder test and the one that speaks to the template-memorisation concern. Because tuning an F1-optimal threshold on the small per-fold partitions is itself a large source of variance, both cross-validation schemes fix $\tau = 0.50$ and report the resulting spread as mean $\pm$ standard deviation over folds, together with the pooled out-of-fold confusion matrix. For the single held-out split we additionally report percentile bootstrap 95% confidence intervals from 2,000 class-stratified resamples [18].

For RQ4, we build an independent holdout of 41 hand-authored real scripts (21 dangerous, 20 safe) drawn from publicly documented idioms—standard sysadmin/devops tasks on the safe side and canonical, publicly catalogued attack techniques (reverse shells, the bash fork bomb, device-wipe dd/mkfs, SSH authorized_keys backdoors, cron persistence, obfuscated payloads) on the dangerous side. None is produced by the synthetic generator and none is byte-identical to any training script. A sequence-similarity check further confirms that no holdout script exceeds 0.85 similarity to any training script (closest pair 0.83). ACTE is trained on the *entire* synthetic corpus (feedback learning and threshold tuning on synthetic data only), then frozen and evaluated once on the real holdout, yielding a genuine train-synthetic / test-real generalization measurement.

VIII. RESULTS AND ANALYSIS

A. RQ1: Detection Accuracy

Table I reports the full metrics for ACTE and each ablation configuration. The full model achieves $F1 = 0.915$, $Precision = 0.985$, $Recall = 0.855$, $MCC = 0.861$, and $ROC-AUC = 0.976$ on the 168-script test set (TN=91, FP=1, FN=11, TP=65). Hypothesis H1 is confirmed. The ROC and precision-recall curves (Figures 1 and 2) sit well above their no-skill baselines, indicating the model ranks scripts by risk effectively; we return to what this does and does not imply in Section IX.

Table I: ACTE Full Model and Ablation Metrics (Test Set, n=168)

Configuration	Prec.	Rec.	F1	MCC	FPR	ROC-AUC
Full ACTE	0.985	0.855	0.915	0.861	0.011	0.976

No SemanticParser	0.957	0.868	0.910	0.846	0.033	0.967
No ContextExtractor	0.985	0.882	0.931	0.883	0.011	0.969
No ThreatIntel	0.833	0.789	0.811	0.663	0.130	0.896
No FeedbackLearning	0.766	0.947	0.847	0.710	0.239	0.952




Fig. 1. ROC curve for the full ACTE model on the held-out test set (n=168). AUC=0.976.




Fig. 2. Precision-Recall curve for the full ACTE model. PR-AUC=0.978.




Fig. 3. F1 scores across ablation configurations. Removing ThreatIntel produces the largest F1 drop (-0.105). Removing FeedbackLearning sharply raises FPR from 0.011 to 0.239.

To check that the single-split headline is not an artifact of one partition, Table I-B reports bootstrap 95% confidence intervals for the held-out metrics alongside the two cross-validation schemes (both at the fixed default $\tau = 0.50$, which is why their F1 sits slightly below the threshold-tuned single-split figure). Stratified five-fold gives $F1 = 0.887 \pm 0.025$, consistent with the single-split result and tight across folds. Leave-template-out five-fold gives $F1 = 0.823 \pm 0.101$: lower and more variable, as one would expect when the model must judge template structures it never trained on, but well clear of chance. We deliberately surface the accompanying false-positive rate, which the headline 0.011 does not represent: at the untuned operating point the stratified FPR is 0.10 ± 0.04 and the leave-template-out FPR is 0.11 ± 0.19 . In other words, the very low headline FPR is a property of the tuned threshold, and under leave-template-out generalisation the false-positive advantage narrows considerably. The honest reading is that leave-template-out cross-validation weakens—without fully dispelling—the concern that a synthetic corpus rewards template memorisation. Figure 5 shows the per-fold F1 distributions.

Table I-B: Robustness of RQ1: bootstrap 95% CIs on the held-out split (tuned τ) and cross-validated metrics (mean $\pm$ std, $k=5$, fixed $\tau = 0.50$)

Metric / Scheme	Precision	Recall	F1	FPR
Held-out point	0.985	0.855	0.915	0.011
Bootstrap 95% CI	[.953, 1.000]	[.776, .934]	[.861, .959]	[.000, .033]
Stratified 5-fold	0.884 ± 0.036	0.895 ± 0.055	0.887 ± 0.025	0.100 ± 0.040
Leave-template-out	0.903 ± 0.142	0.800 ± 0.175	0.823 ± 0.101	0.113 ± 0.186




Fig. 5. Per-fold F1 under stratified vs. leave-template-out five-fold cross-validation. The stratified scheme is tight around 0.90; the leave-template-out scheme is lower and more dispersed, quantifying the harder task of generalizing to unseen script structures.

B. Ablation Analysis

The ablation results confirm that ThreatIntel is the single most discriminative component: its removal reduces F1 by 0.105 and increases FPR from 0.011 to 0.130. FeedbackLearning, while contributing a smaller F1 change ($\Delta F1 = -0.068$), is the primary mechanism controlling false positives: without it, FPR rises dramatically from 0.011 to 0.239 as the unlearned model defaults to over-predicting dangerous. Removing ContextExtractor, counter-intuitively, yields a marginally higher F1 (0.931) due to threshold shift; this suggests that in the present dataset, context features add discriminative precision at the cost of some recall, a trade-off resolvable via threshold tuning. SemanticParser removal has the smallest individual effect ($\Delta F1 = -0.005$) because the remaining components partially compensate through overlapping signal.

C. RQ2: Computational Cost

Mean per-script analysis latency is 0.52 ms (median 0.48 ms, p95 0.96 ms, p99 1.26 ms), measured over 420 scripts with three repetitions each and the minimum kept per script to suppress scheduler jitter; throughput is roughly 1,900 scripts/second. Latency is the one figure that varies run to run, since it is a wall-clock measurement. Hypothesis H2 holds comfortably: even the 95th-percentile stays under a millisecond, well inside the budget for an inline CI/CD or shell-hook check. These numbers cover analysis only; installing a live seccomp filter would add a one-time per-execution cost that the prototype does not incur.

D. RQ3: Comparison with a Static Linter

Table II compares ACTE against ShellCheck v0.9.0. The point of this comparison is not to claim a scalp—ShellCheck is a style and correctness linter, and using it as a danger

classifier is a category error—but to make concrete why teams cannot lean on the tool they already run. Under the charitable error+warning mapping (mean 0.27 findings/script) ShellCheck flags 17.4% of safe scripts while catching under 8% of dangerous ones (FPR 0.174, recall 0.079, F1 0.122, MCC −0.140); the negative MCC means its "dangerous" verdicts are, if anything, slightly anti-correlated with actual danger. The stricter error-only mapping never fires on a safe script but catches a single true positive (recall 0.013). ACTE's FPR of 0.011 is far lower, but the honest way to state the result is qualitative: a linter and a trust classifier are answering different questions, and RQ5 provides the comparison that actually tests ACTE's modelling choices.

Table II: ACTE vs. ShellCheck v0.9.0 Baseline Comparison (Test Set)

System	Prec.	Recall	F1	MCC	FPR
ACTE (full)	0.985	0.855	0.915	0.861	0.011
ShellCheck (error+warn)	0.273	0.079	0.122	−0.140	0.174
ShellCheck (error only)	1.000	0.013	0.026	0.085	0.000




Fig. 4. ACTE vs. ShellCheck v0.9.0 across precision, recall, and FPR on the test set. ShellCheck, a linter rather than a security tool, both misses most dangerous scripts and raises more false alarms.

The gap is not a sampling accident: the exact McNemar test on the paired predictions finds ACTE correct on 75 items where ShellCheck errs against one the other way, $p \approx 2 \times 10^{-21}$. We read this as confirming that a linter is the wrong instrument for the job, not as evidence that ACTE is a strong detector in absolute terms—that claim rests on RQ5.

E. RQ4: Generalization to Real Scripts

Table III reports the external-validation result. Trained only on the synthetic corpus, frozen, and run once on the 41-script holdout, ACTE reaches F1 = 0.895 (precision 1.000, recall 0.810, ROC-AUC 0.990, FPR 0.000; F1 bootstrap 95% CI [0.79, 0.97], necessarily wide at this sample size). The precision is the part we find most encouraging: all 20

safe scripts are cleared, including the traps we planted—a recursive `chmod` on a web root, a `dd` copy between two regular files, a privileged `apt` install—while the trusted-vendor `curl | sh` installer is flagged, consistent with our conservative stance that piping remote code into a shell is the hazard regardless of the vendor. The four misses are the honest part. All are evasive techniques with no matching signature and only moderate context scores: a `socat` reverse shell, a `blkdiscard` device wipe, an `rsync --delete` that empties the filesystem without ever naming `rm`, and a ROT13-obfuscated download-and-run. The high ROC-AUC (0.990) alongside the lower recall tells us the ranking is sound but the fixed threshold is conservative on unfamiliar payloads—precisely the case for the signature-induction and behavioural-feature work we flag as future work, and a reminder that a finite signature base is a coverage floor, not a ceiling.

Table III: Real-World External Validation (train on synthetic, test on 41 real scripts)

Metric	Precision	Recall	F1	ROC-AUC	FPR
ACTE (frozen)	1.000	0.810	0.895	0.990	0.000




Fig. 6. Metrics on the real-world holdout for a model trained only on synthetic data: zero false positives on 20 real safe scripts, recall 0.810 with four evasive payloads missed, ROC-AUC 0.990.

F. RQ5: Against Learned Baselines

The comparison that actually stresses ACTE's design is against classifiers trained on the same labels. We fit TF-IDF features (word 1–2 grams and character 3–5 grams) to logistic regression, a linear SVM, and a random forest. Table IV shows the outcome, and it is not the outcome a paper usually reports about its own method: on the synthetic test split the linear text classifiers reach F1 = 0.974, above ACTE's 0.915, and on the real holdout the linear SVM reaches F1 = 0.977 to ACTE's 0.895. We state this plainly because it sharpens what ACTE is actually for. A thirteen-feature logistic model does not win a raw-accuracy contest

against a few thousand lexical features, and it was never going to.

What ACTE does hold is the properties that matter for a gate. Its false-positive rate is the lowest of the field on the real holdout (0.000 against 0.050 for the linear models and 0.300 for the random forest), and on a gate the cost of a false alarm is what operators actually pay. Its decision is attributable to thirteen named features with signed contributions, whereas the logistic baseline’s verdict rests on opaque weights over tokens such as the bare string `rf` and on distributional accidents of the corpus (the synthetic `.test` hostnames earn their own weight); such a model is both hard to audit and easy to evade by renaming. Its weights update online from a single labelled example in under a millisecond, where a fitted TF-IDF vocabulary is frozen at training time. And it emits an enforcement policy, which a bare label cannot. The reasonable conclusion is not that ACTE beats these baselines but that it occupies a different, arguably more deployable point in the design space: interpretable, adaptive, policy-producing, and calibrated toward precision.

Table IV: ACTE vs. learned TF-IDF baselines. Left: synthetic test split. Right: real-world holdout (trained on the full synthetic corpus).

Detector	F1 (test)	FPR (test)	F1 (real)	FPR (real)
ACTE (full)	0.915	0.011	0.895	0.000
TF-IDF + LogReg	0.974	0.022	0.952	0.050
TF-IDF + LinearSVM	0.974	0.022	0.977	0.050
TF-IDF + RandomForest	0.968	0.054	0.875	0.300




Fig. 7. F1 and FPR on the real-world holdout. The learned text classifiers match or beat ACTE on F1 but at a higher false-positive rate; ACTE alone reaches zero false positives.

G. RQ6: Robustness to Adaptive Evasion

RQ5 leaves an uncomfortable question: if a TF-IDF classifier matches ACTE on F1, why prefer ACTE? The

answer is what happens under an adversary. We take the dangerous test scripts and apply two behaviour-preserving edits, leaving the malicious commands untouched. *Benign camouflage* prepends a shebang, set `-euo pipefail`, and a few reassuring comments — a direct probe of the monotonicity property, since those tokens carry negative weight in our model. *Lexical disguise* renames attacker hostnames and variables. Table V reports recall before and after.

Table V: Recall on dangerous scripts before and after behaviour-preserving camouflage

Detector	Original	Benign camouflage	Lexical disguise	Both
ACTE (full)	0.855	0.842	0.855	0.842
TF-IDF + LogReg	0.974	0.316	0.974	0.303

The result reverses RQ5’s raw-F1 verdict. Benign camouflage costs ACTE 0.013 recall; it costs the lexical baseline 0.658 — a two-line, behaviour-preserving edit that any operator would wave through collapses the bag-of-tokens model from 0.974 to 0.316, while ACTE barely moves. The reason is exactly the structure the two models do not share: ACTE’s signature and context weights for the intact malicious commands dominate the small negative weights of the benign-signal tokens, whereas the TF-IDF model had learned that shebangs, set `-e`, and comments mark safe scripts. Renaming hostnames moves neither detector, because both key on command structure, not specific strings. This is also the empirical answer to the monotonicity objection of Section IV-D: the property is a theoretical evasion lever, but the benign-signal weights are too small to float a genuinely dangerous script under the threshold. The evasions that do succeed against ACTE are the ones from RQ4 — novel techniques with no signature — not cosmetic disguises.

H. RQ6b: Systematic Semantic-Substitution Attacks

RQ6 tested additive camouflage; the attacks that actually work are semantic substitutions that swap a signature-matched construct for a behaviour-equivalent one the signature base does not cover. We apply five such transforms to the dangerous test scripts against ACTE and every learned baseline, preserving the malicious effect, and report recall as a function of attacker budget. Table VII places both attack families side by side. The result reverses RQ6: at maximum budget, semantic substitution costs ACTE 0.276 recall (0.855 → 0.579) but costs the lexical baseline nothing, just as additive camouflage cost the lexical baseline 0.658 but

ACTE almost nothing. The two families fail under opposite edits—ACTE when its signatures are routed around, the lexical model when benign tokens are added—so robustness is not a reason to prefer one; it is a reason to combine structural and lexical evidence, which the union model of RQ8 begins to do.

Table VII: Recall of ACTE vs. the lexical baseline under two behaviour-preserving attack families (dangerous test scripts)

Detector	Original	+ Additive camouflage (RQ6)	+ Semantic substitution (RQ6b)
ACTE	0.855	0.842	0.579
TF-IDF + LogReg	0.974	0.316	0.974

I. RQ7: Calibration of the Risk Score

The four trust levels assume the score behaves like a probability. It only roughly does: on the test split the raw Expected Calibration Error is 0.076 (Brier 0.061), so the score is systematically over- or under-confident by about eight points. Isotonic recalibration fit on the training split reduces ECE to 0.055; Platt scaling does not help (0.084). We therefore qualify the “graded enforcement proportionate to risk” language: the tier boundaries (0.25 / 0.50 / 0.80) are reasonable but not calibrated, and a deployment that maps tiers to costly actions should recalibrate on its own data [30]. This supplies the evidence the earlier draft asserted without.

J. RQ8: Baselines that Isolate Features from the Model

Comparing a hand-tuned logistic model to TF-IDF conflates the feature set with the linear form, so we add three sharper baselines (Table VI). A gradient-boosted tree on the *same* thirteen features reaches F1 0.923—marginally above the logistic ACTE—so the ceiling on this corpus is the features, not linearity. Concatenating the thirteen features with TF-IDF reaches F1 0.993 (ROC-AUC 1.000), which shows both that the engineered features add signal on top of lexical ones and, candidly, that the synthetic corpus is close to separable once enough features are combined—a limitation of the benchmark, not a triumph of the model. Tuning the TF-IDF baseline’s regularisation (grid over C, train-only) leaves it at F1 0.974, so RQ5 was not a comparison against a crippled baseline. A frozen code-transformer baseline (CodeBERT/StarEncoder) would sharpen the picture further, but its weights cannot be fetched in our offline environment; we flag it as future work rather than fabricate it.

Table VI: Stronger baselines on the synthetic test split, isolating features from the linear model

Detector	F1	FPR	ROC-AUC
----------	----	-----	---------

ACTE (logistic, 13 feat)	0.915	0.011	0.976
GBDT (same 13 feat)	0.923	0.011	0.972
Feature + TF-IDF union	0.993	0.011	1.000
TF-IDF + LogReg (tuned, C=4)	0.974	0.022	0.999

K. RQ9: Does Online Learning Earn Its Place?

Two experiments test the online loop the paper originally sold. For concept drift we hold out the obfuscated family as a distribution the model has not trained on; the frozen model already detects it at full recall, so online adaptation recovers nothing. Adapting instead on the dangerous test scripts the frozen model genuinely misses does recover them—but drives the false-positive rate on the safe test set to 1.0, i.e. it overfits catastrophically. For poisoning we feed a genuinely dangerous canary’s own features with a “safe” label: its risk falls 0.999→0.382 and crosses the threshold after six updates. Online learning, as implemented, is thus neither demonstrably beneficial nor safe against label poisoning, and we no longer present it as a contribution (Section X-E).

L. RQ10: Real, Content-Derived seccomp Enforcement

To close the construct-validity gap that “runtime security” was never enforced, we derive, from the commands ACTE parses out of a script, the syscalls it legitimately needs and hence a per-script *deny*-set (a script with no network command may not `socket/connect`; one that spawns nothing may not `execve`; none may `ptrace`). We compile that deny-set to a classic-BPF seccomp filter and install it in a forked child, unprivileged, via `PR_SET_NO_NEW_PRIVS`. Across five probe scripts the kernel enforced the policy correctly every time—a denied syscall was killed with `SIGSYS` and a control syscall permitted—and two scripts with different commands received different filters, i.e. genuine content-derived synthesis, really enforced, rather than one of a few per-tier templates. This is a demonstrative least-privilege gate, not a full sandbox (Section X-C), but it moves the enforcement claim from aspiration to a working kernel control.

IX. STATISTICAL TREATMENT

Point estimates on a single split cannot, by themselves, separate a real effect from a lucky partition, so we quantify uncertainty three ways. Percentile bootstrap 95% intervals from 2,000 class-stratified resamples of the 168-sample test set give F1 [0.861, 0.959], precision [0.953, 1.000], recall [0.776, 0.934], and ROC-AUC [0.951, 0.994] (Table I-B). Because resampling holds class proportions fixed, these intervals understate uncertainty from class-balance sampling and, on the small holdouts, are wide; we

present them as bounds, not as tight guarantees. Five-fold cross-validation puts F1 at 0.887 ± 0.025 (stratified) and 0.823 ± 0.101 (leave-template-out), so the result reproduces under resampling and largely survives the removal of template overlap. The exact McNemar test rejects equal error rates against ShellCheck at $p \approx 2 \times 10^{-21}$; as discussed in Section VIII-D we treat this as evidence about linters, not about competing detectors.

The Matthews Correlation Coefficient (MCC) of 0.861 is an informative summary statistic for imbalanced binary classification, as it takes into account all four entries of the confusion matrix and is not inflated by class imbalance. An MCC approaching 1.0 indicates near-perfect classification; our value of 0.861 is consistent with a high-performing but not perfect detector, with the primary source of imperfection being false negatives in the malicious category (recall=0.750).

The ROC and PR curves (Figures 1 and 2) show that the ranking quality is threshold-independent: ROC-AUC 0.976 and PR-AUC 0.978 mean the model orders scripts well irrespective of where the cut is placed. We are careful not to call the model "calibrated": ranking quality (AUC) is a separate property from calibration (whether a risk of 0.7 corresponds to a 70% empirical danger rate), and we have not measured the latter. A reliability analysis, and threshold selection that accounts for it, is left to future work.

Finally, we note a limitation of the ablation deltas: they are point differences between configurations trained under one protocol on one split, not paired significance tests, and the confidence intervals of neighbouring configurations overlap. They indicate the direction and rough magnitude of each component's contribution, not a certified ranking.

Two corrections to the statistical treatment above. First, the bootstrap intervals in Table I-B resampled the 168 test scripts as if independent, but scripts are clustered by generating template, so those intervals were too narrow. A cluster (block) bootstrap that resamples whole templates gives an F1 interval of [0.79, 0.99] on the held-out split—about twice the width of the naive interval—with comparable widening for recall; the honest headline uncertainty is therefore substantially larger than first reported. Second, the “zero false positives” on the real-world holdout is 0 of 20 negatives, whose Wilson 95% interval is [0.00, 0.16]—consistent with a true false-positive rate as high as ~16%, so it is a point estimate on a small sample, not a guarantee.

The ablation deltas are now backed by paired McNemar tests on the same test items rather than read from point estimates. Only two components differ significantly from the full model: removing ThreatIntel ($p = 0.007$) and removing FeedbackLearning ($p = 0.013$). Removing SemanticParser ($p = 1.0$) or ContextExtractor ($p = 0.5$) does not, so we withdraw any ranking of those two components—including the earlier observation that removing ContextExtractor “improves” F1, which is within noise.

X. THREATS TO VALIDITY

A. Internal Validity: a closed loop

The most serious threat to this study is that one author wrote the threat signatures, the feature extractors, the synthetic-corpus generator, and the real-world holdout. When the designer of a detector also authors the data it is judged on, the evaluation risks measuring self-consistency rather than detection ability, and the ablation makes the dependence explicit: removing ThreatIntel is the single largest hit to F1, which is to say a good deal of the performance rests on author-written regexes matching payloads drawn from an author-known vocabulary of techniques. Leave-template-out cross-validation and the independence-checked real holdout push against this concern but cannot dissolve it. We therefore read the absolute numbers as evidence about the *method* under controlled conditions, and we do not present them as an estimate of field detection rate. The decisive test—third-party scripts (real assistant outputs, public malware-shell corpora) labelled by independent annotators with reported inter-annotator agreement—is exactly the study we cannot run from author-generated data, and it is the first item of future work for a reason.

B. External Validity: synthetic data and its confounds

The training corpus is synthetic: 420 scripts from 70 templates, so the effective number of independent items is closer to 70 than to 420, and category is largely confounded with label (the *malicious* and *obfuscated* categories are entirely dangerous, *safe_everyday* entirely safe, and the *sysadmin* category contains a single dangerous script, whose per-category “recall = 1.0” is therefore anecdotal). The honest per-category picture is that recall on the genuinely hard *malicious* class is 0.75, and the abstract's 0.855 aggregate is buoyed by easier categories. We mitigate the memorisation concern with leave-template-out cross-validation (F1 = 0.823) and the transfer concern with the RQ4 holdout (F1 = 0.895, precision = 1.000), the latter verified to share no script above 0.83 similarity with training. But 41 real scripts is a small, author-curated

sample; its confidence intervals are wide and it is not a substitute for field data. None of the absolute numbers should be assumed to transfer unchanged to production.

C. Construct Validity: enforcement is now real, but partial

Enforcement is no longer only a sketch. Section VIII-L compiles a content-derived syscall deny-set into an actual classic-BPF seccomp filter, installs it unprivileged via `PR_SET_NO_NEW_PRIVS`, and shows the Linux kernel kill a denied system call with `SIGSYS` while permitting the controls—correctly on all five probe scripts, with a different filter per script. That validates the score $\rightarrow$ policy $\rightarrow$ kernel-control path end to end at the system-call boundary the threat model names. What remains a sketch is the surrounding runtime: we do not execute the scripts inside a managed container with the full namespace/cgroup profile applied, and the OPA/Rego and cgroup artefacts are still emitted as inspectable data rather than mounted. So “runtime security” is now demonstrated at the seccomp layer but not yet as a complete sandbox; integrating the artefacts into a container runtime (containerd/Podman) with the whole profile enforced is the remaining systems work.

D. Adversarial Robustness: complementary blind spots

Sections VIII-G and VIII-H probe the adversary from two directions, and the honest result is that neither ACTE nor a lexical model is robustly superior. Additive benign camouflage—the direct exploit of monotonicity—barely dents ACTE (0.013 recall) but collapses the TF-IDF baseline (recall 0.97 $\rightarrow$ 0.32), consistent with monotone classifiers being provably robust to additive, benign feature injection [26]. Behaviour-preserving *semantic substitution*—replacing a signature-matched construct with an equivalent one (`rm -rf` $\rightarrow$ `find -delete`, `nc` $\rightarrow$ `-e socat, curl|bash` $\rightarrow$ `staged download-and-exec`)—does the reverse: at maximum attacker budget it evades ACTE (recall 0.86 $\rightarrow$ 0.58) while leaving the lexical baseline unaffected. This is exactly the problem-space versus feature-space distinction of Pierazzi et al. [23]: monotonicity defends only against additive problem-space edits, and the real weakness is signature/feature coverage, which an attacker who knows the finite signature base can route around. We therefore reframe monotonicity (Section IV-D) as a transparency property with a security cost, not a robustness guarantee, and we make no claim of adversarial robustness. RQ6b is a first systematic step; a full adaptive-attacker study with an explicit budget against every detector, including attacks tailored to each, remains future work.

E. The Online-Learning Loop

We include online SGD adaptation but do not validate it as a benefit. On a held-out “novel” family the frozen model already reaches full recall, so adaptation adds nothing; and adapting on the model’s genuine misses recovers them only by driving the false-positive rate to 1.0 (Section VIII-K). The loop is also an attack surface: feeding a dangerous script’s own feature vector with a “safe” label six times flips its decision. A deployable feedback loop would need authenticated, rate-limited, multi-party labels and regularised or replay-based updates; we treat online learning as future work, not a result, and have removed it from the paper’s claimed contributions.

XI. REPRODUCIBILITY

The code, data generators, experiment runners, and figure scripts are in the public repository. One command, `python -m experiments.run_all`, regenerates the synthetic corpus (seed 1337) and the real-world holdout, trains the model, runs all five research questions—the two cross-validation schemes, bootstrap intervals, the McNemar test, the learned baselines, and the external validation—and writes `results.json`, `results.md`, and every figure. Corpus membership is now content-hashed with SHA-256 rather than the salted built-in `hash()`, so the dataset is byte-identical across runs and interpreters; the grouped folds use scikit-learn’s `StratifiedGroupKFold`. A suite of over 140 tests (`python -m pytest tests/`) guards every component and encodes the invariants a reviewer would want checked: that leave-template-out folds share no template, that fold class balance is not degenerate, and that no holdout script exceeds 0.85 sequence similarity to any training script.

Provenance is documented in `data/PROVENANCE.md` for the synthetic corpus and in `data/real_world/build.py` for the holdout, the latter recording each script’s public source and its labelling rationale. Python 3.11 on Linux was used for the reported run; dependencies are pinned in `requirements.txt`. Every classification metric is reproducible given the same seed; only wall-clock latency varies by design.

XII. RELATED WORK

A. Static Analysis and Linting

ShellCheck [1] is the most widely deployed static analyzer for shell scripts, detecting syntactic bugs, quoting errors, and portability issues via a combination of grammar-based parsing and heuristic rules. Its scope is correctness and style

rather than security, and it operates entirely statically without a concept of trust or adaptive scoring. Our evaluation quantifies the practical limitation of applying ShellCheck as a security filter: the error+warning mapping achieves recall of only 7.9% at FPR of 17.4% on our test set, confirming its unsuitability for the threat model considered here.

B. Mandatory Access Control

SELinux [2] and AppArmor [3] are kernel-enforced mandatory access-control frameworks. SELinux implements a type-enforcement model in which every process and object is labeled and access decisions follow an explicit policy database. AppArmor uses pathname-based profiles. Both are highly effective at constraining known applications, but they require human-authored per-process policies; they cannot automatically generate an appropriate policy from the content of an unseen script at submission time. ACTE's PolicyGenerator addresses this gap by synthesizing enforcement artifacts directly from the risk assessment.

C. Behavioral Monitoring

Falco [4] is a runtime threat-detection tool that monitors kernel-level system calls and container events against a rule set, alerting on anomalous behavior. Linux auditd [5] provides a similar audit log at the syscall level. eBPF-based monitors extend this capability with programmable, low-overhead kernel probes [6]. These systems are inherently *reactive*: they detect dangerous behavior after it begins. ACTE is *proactive*: it assigns a trust level before execution and generates an enforcement policy that constrains what the script is permitted to do, rather than logging what it did.

D. Sandboxing and Policy Enforcement

seccomp-bpf [7] provides a kernel-enforced syscall filter that can deny or kill a process on forbidden system calls. Linux namespaces [8] and cgroups [9] isolate process view of the system and constrain resource consumption respectively. Open Policy Agent (OPA) and its Rego policy language [10] provide a general-purpose, declarative policy engine used in cloud-native stacks. ACTE builds on this ecosystem: its PolicyGenerator emits seccomp profiles and Rego policies, and its RuntimeMonitor simulates seccomp-bpf enforcement. The novelty lies in the automated derivation of these policies from a content-based trust score rather than requiring manual authorship.

E. LLM-Generated Code Security

A growing body of work studies the security properties of LLM-generated code. Pearce et al. [11] found that GitHub

Copilot introduces security vulnerabilities in approximately 40% of generated code samples across a range of CWE categories. Khoury et al. [12] evaluated ChatGPT-generated code for security vulnerabilities, finding frequent occurrences of SQL injection, path traversal, and command injection. Tony et al. [13] introduced LLMSecEval, a benchmark for evaluating LLMs on security-sensitive code generation tasks. These works focus on code quality and vulnerability introduction at generation time; ACTE complements them by providing a runtime gate that intercepts potentially dangerous scripts regardless of how they were generated.

F. Zero-Trust Execution Models

The zero-trust security paradigm [14] posits that no entity should be inherently trusted, and access decisions should be continuously evaluated based on context and evidence. ACTE instantiates a zero-trust approach specifically for code execution: every submitted script receives a trust evaluation from first principles, and execution permissions are granted proportionate to the assessed trust level rather than based on the submitter's identity or the script's syntactic validity. This aligns with emerging work on continuous authorization for automated systems [15].

G. Learning-Based Detection of Obfuscated Shell and PowerShell

Learning-based detection of obfuscated scripting attacks has developed along two representational axes. Token- and character-level models treat a command as a sequence: Hendler, Kels, and Rubin combined NLP-style detectors with a character-level convolutional network to flag malicious PowerShell command lines, showing that lexical patterns often survive obfuscation transforms [19]. Structural approaches instead parse the script before classifying. Rusak, Al-Dujaili, and O'Reilly learned embeddings over abstract-syntax-tree (AST) nodes, arguing that syntactic structure is harder for an adversary to perturb than raw text [20]. Because PowerShell can construct code fragments at runtime, Li et al. paired subtree-level deobfuscation with semantic-signature mining, recovering roughly 80% of the original script similarity before classification and distilling 31 reusable attack signatures [21]. Most relevant to shell scripts specifically, SCORE generalizes syntactic representation beyond PowerShell, extracting both syntax-highlighting token sequences (SCORE-H) and AST features (SCORE-T) for static script-malware detection across multiple languages [22]. Together these results motivate fusing token- and AST-derived

signals, yet they largely evaluate against a static, non-adaptive adversary—precisely the assumption a signature-informed detector deployed against AI-generated scripts must revisit.

H. Adversarial ML for Malware: Problem-Space vs. Feature-Space

A central lesson of adversarial ML for malware is that gradient-space evasion and real-world evasion are not the same thing. Pierazzi et al. formalize this as the gap between feature-space perturbations and problem-space attacks, which must produce a real, functional artifact under available-transformation, preserved-semantics, robustness-to-preprocessing, and plausibility constraints, and they introduce the notion of side-effect features generated by any concrete edit [23]. Concrete problem-space attacks bear this out: Kolosnjaji et al. evade a byte-level network by appending less than 1% padding while preserving functionality [24], and Demetrio et al. demonstrate query-efficient, functionality-preserving black-box evasion by injecting never-executed benign content [25]. This distinction directly shapes defense design. Because an attacker can freely add content but cannot cheaply remove the malicious core, Incer et al. advocate monotonic classifiers whose score can only increase as features are added, so evasion cannot be achieved by appending benign material [26]. A monotone, signature-informed classifier inherits exactly this property: since obfuscation and padding are additive problem-space transforms that introduce side-effect features rather than deleting signatures, monotonicity forces the adversary into the far costlier regime of removing genuinely malicious behavior.

I. Explainable Malware Detection

Interpretability is increasingly treated not as a post-hoc nicety but as a first-class requirement for security classifiers, both to earn analyst trust and to expose brittle decision logic. Guo et al.'s LEMNA tailors local explanations to security applications, handling the feature dependence and nonlinear decision boundaries that generic image-oriented explainers violate, and it shows explanations being used to validate and even patch malware-classifier behavior [27]. Warnecke et al. systematize this line of work by defining evaluation criteria—descriptive accuracy, sparsity, stability, robustness, and efficiency—and comparing six explanation methods on security tasks including malware detection, giving concrete guidance on which explainers are trustworthy in this domain [28]. Explanations also double as an attack-surface diagnostic: Demetrio et al. apply integrated gradients to the

MalConv detector and reveal that it keys on PE-header bytes rather than malicious content, an interpretability finding that directly predicts a header-only evasion [29]. Taken together, these works support treating interpretability as a feature of the detector itself: a signature-informed model whose decisions decompose into human-auditable, named indicators lets analysts confirm that a flag rests on genuinely malicious constructs rather than spurious correlations, aligning explanation fidelity with the robustness goals above.

XIII. CONCLUSION AND FUTURE WORK

ACTE scores a shell script before it runs, through a transparent thirteen-feature logistic model over semantic, execution-context, and threat-signature evidence, and compiles the verdict into an enforcement policy that—on Linux—we install as a real seccomp filter and show the kernel enforce end to end. We have tried to report the work as a reviewer would want it read: the honest generalisation figure is a leave-template-out F1 of 0.823 ± 0.101 (cluster-bootstrap F1 interval [0.79, 0.99]), recall on the hard malicious class is 0.75, and the default-threshold false-positive rate is about 0.10, not the 0.011 of one tuned split. Against strong learned baselines ACTE is competitive, not dominant, on raw F1; its value is the lowest false-positive rate, an auditable decision, and a really-enforced content-derived policy. Under adaptive evasion it and a lexical model have complementary blind spots, so neither is robust. We do not claim online learning as a benefit (it is unvalidated and poisonable), and we are explicit that the training corpus is synthetic and author-generated and that a full sandbox and independent human labelling remain undone. We offer ACTE as a reproducible, honestly-evaluated framework for interpretable content-to-policy trust gating, not as a finished detector.

This paper is the first in a planned research program. Immediate future work includes: (1) scaling the external validation of Section VIII-E into a large-scale field study that harvests AI-generated scripts from deployed LLM coding assistants with human-expert labeling, extending the promising 41-script generalization result; (2) integration of ACTE with a live container runtime (containerd or Podman) to deploy the generated seccomp profiles and cgroup constraints as actual kernel-level enforcement rather than sketches; (3) extension of the threat-intelligence knowledge base with community-sourced signatures and automated signature induction from observed incidents; (4) a formal analysis of the monotonicity property and its implications

for adversarial robustness under feature-space attacks; and (5) evaluation of ACTE in a multi-tenant CI/CD environment to characterize operational false-positive burden and operator feedback loop latency.

REFERENCES

1. V. Holen, "ShellCheck: A static analysis tool for shell scripts," *Software*, 2012. [Online]. Available: <https://www.shellcheck.net>
2. P. Loscocco and S. Smalley, "Integrating flexible support for security policies into the Linux operating system," in *Proc. USENIX Annual Technical Conference (FREENIX Track)*, 2001, pp. 29–42.
3. AppArmor Security Project, "AppArmor application security," project documentation, 2024. [Online]. Available: <https://apparmor.net>
4. Falco Authors, "Falco: Cloud-native runtime security," *CNCF Project*, 2016. [Online]. Available: <https://falco.org>
5. The Linux Audit Project, "Linux audit (auditd) documentation." [Online]. Available: <https://github.com/linux-audit/audit-documentation>
6. B. Gregg, *BPF Performance Tools: Linux System and Application Observability*. Addison-Wesley, 2019.
7. The Linux Kernel, "Seccomp BPF (SECure COMPUting with filters)," kernel documentation. [Online]. Available: https://www.kernel.org/doc/html/latest/userspace-api/seccomp_filter.html
8. M. Kerrisk, "Namespaces in operation," *LWN.net*, 2013. [Online]. Available: <https://lwn.net/Articles/531114/>
9. P. Menage, "Adding generic process containers to the Linux kernel," in *Proc. Linux Symposium*, 2007, pp. 45–57.
10. Open Policy Agent Authors, "Open Policy Agent: Policy-based control for cloud native environments," *CNCF Project*, 2016. [Online]. Available: <https://www.openpolicyagent.org>
11. H. Pearce, B. Ahmad, B. Tan, B. Dolan-Gavitt, and R. Karri, "Asleep at the keyboard? Assessing the security of GitHub Copilot's code contributions," in *Proc. IEEE Symposium on Security and Privacy (S&P)*, 2022, pp. 754–768.
12. R. Khoury, A. R. Avila, J. Brunelle, and B. M. Camara, "How secure is code generated by ChatGPT?" in *Proc. IEEE International Symposium on Systems Engineering*, 2023.
13. C. Tony, M. Mutas, N. E. Díaz Ferreyra, and R. Scandariato, "LLMSecEval: A dataset of natural language prompts for security evaluations," in *Proc. IEEE/ACM Mining Software Repositories (MSR)*, 2023.
14. NIST, "Zero trust architecture," *NIST Special Publication 800-207*, National Institute of Standards and Technology, 2020.
15. J. Kindervag, "Build security into your network's DNA: The zero trust network architecture," *Forrester Research*, Tech. Rep., 2010.
16. F. Pedregosa *et al.*, "Scikit-learn: Machine learning in Python," *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
17. T. G. Dietterich, "Approximate statistical tests for comparing supervised classification learning algorithms," *Neural Computation*, vol. 10, no. 7, pp. 1895–1923, 1998.
18. B. Efron and R. J. Tibshirani, *An Introduction to the Bootstrap*. Chapman & Hall/CRC, 1993.
19. D. Hendler, S. Kels, and A. Rubin, "Detecting Malicious PowerShell Commands using Deep Neural Networks," in *Proc. ACM ASIACCS*, 2018, pp. 187–197.
20. G. Rusak, A. Al-Dujaili, and U.-M. O'Reilly, "POSTER: AST-Based Deep Learning for Detecting Malicious PowerShell," in *Proc. ACM CCS*, 2018, pp. 2276–2278.
21. Z. Li, Q. A. Chen, C. Xiong, Y. Chen, T. Zhu, and H. Yang, "Effective and Light-Weight Deobfuscation and Semantic-Aware Attack Detection for PowerShell Scripts," in *Proc. ACM CCS*, 2019, pp. 1831–1847.
22. E. Erdemir, K. Park, M. J. Morais, V. R. Gao, M. Marschalek, and Y. Fan, "SCORE: Syntactic Code Representations for Static Script Malware Detection," *arXiv:2411.08182*, 2024.
23. F. Pierazzi, F. Pendlebury, J. Cortellazzi, and L. Cavallaro, "Intriguing Properties of Adversarial ML Attacks in the Problem Space," in *Proc. IEEE S&P*, 2020, pp. 1332–1349.
24. B. Kolosnjaji *et al.*, "Adversarial Malware Binaries: Evading Deep Learning for Malware Detection in Executables," in *Proc. EUSIPCO*, 2018, pp. 533–537.
25. L. Demetrio, B. Biggio, G. Lagorio, F. Roli, and A. Armando, "Functionality-Preserving Black-Box Optimization of Adversarial Windows Malware," *IEEE Trans. Information Forensics and Security*, vol. 16, pp. 3469–3478, 2021.
26. I. Incer, M. Theodorides, S. Afroz, and D. Wagner, "Adversarially Robust Malware Detection Using Monotonic Classification," in *Proc. ACM IWSPA*, 2018, pp. 54–63.
27. W. Guo, D. Mu, J. Xu, P. Su, G. Wang, and X. Xing, "LEMNA: Explaining Deep Learning based Security Applications," in *Proc. ACM CCS*, 2018, pp. 364–379.
28. A. Warnecke, D. Arp, C. Wressnegger, and K. Rieck, "Evaluating Explanation Methods for Deep Learning in Security," in *Proc. IEEE EuroS&P*, 2020, pp. 158–174.
29. L. Demetrio, B. Biggio, G. Lagorio, F. Roli, and A. Armando, "Explaining Vulnerabilities of Deep Learning to Adversarial Malware Binaries," in *Proc. ITASEC*, CEUR-WS vol. 2315, 2019.
30. C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, "On Calibration of Modern Neural Networks," in *Proc. ICML*, 2017, pp. 1321–1330.